

s-in-nlp-using-bilstm-kikuyu-words

March 16, 2024

```
[91]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from keras.preprocessing.sequence import pad_sequences
from keras.utils import to_categorical
from keras.models import Sequential
from keras.layers import Embedding, Bidirectional, LSTM, Dense, Flatten
from sklearn.preprocessing import OneHotEncoder
from keras.utils import to_categorical
from sklearn.preprocessing import LabelEncoder
```

```
[92]: # Load the dataset
data = pd.read_csv('/content/Kikuyu_Words.csv')
```

```
[93]: # Viewing the first 25 words with corresponding POS label
data.head(25)
```

```
[93]:
```

	Word	Label
0	Mündũ	Noun
1	Mũaki	Noun
2	Ihũa	Noun
3	Kĩrĩma	Noun
4	Gĩtĩri	Noun
5	Ikara	Verb
6	Rehe	Verb
7	Tengera	Verb
8	Rũga	Verb
9	koma	Verb
10	andika	Verb
11	tuma	Verb
12	Cukuru	Noun
13	Handũ	Noun
14	Mūteti	Noun
15	Mũiko	Noun
16	Mũitu	Noun
17	Ndereba	Noun
18	Mũrogi	Noun

```

19      Kĩhĩĩ  Noun
20  Mũrutani  Noun
21      Mũirũ  Noun
22      Mũici  Noun
23      Kiratũ  Noun
24      Mũrũthi  Noun

```

```

[94]: # Preprocessing for easy tokenization
data["Word"] = data["Word"].str.replace("ũ", "u")
data["Word"] = data["Word"].str.replace("ĩ", "i")
data['Word'] = data['Word'].str.lower()

```

```

[95]: data.head(25)

```

```

[95]:      Word Label
0      mundu  Noun
1      muaki  Noun
2      ihua  Noun
3      kirima  Noun
4      gitiri  Noun
5      ikara  Verb
6      rehe  Verb
7      tengera  Verb
8      ruga  Verb
9      koma  Verb
10     andika  Verb
11     tuma  Verb
12     cukuru  Noun
13     handu  Noun
14     muteti  Noun
15     muiko  Noun
16     muitu  Noun
17     ndereba  Noun
18     murogi  Noun
19     kihii  Noun
20  murutani  Noun
21     muiru  Noun
22     muici  Noun
23     kiratu  Noun
24     muruthi  Noun

```

```

[96]: # Encoding labels
label_encoder = LabelEncoder()
data['Label'] = label_encoder.fit_transform(data['Label'])

```

```

[97]: # viewing words to be used as features
X = data['Word'].values

```

X

```
[97]: array(['mundu', 'muaki', 'ihua', 'kirima', 'gitiri', 'ikara', 'rehe',  
          'tengera', 'ruga', 'koma', 'andika', 'tuma', 'cukuru', 'handu',  
          'muteti', 'muiko', 'muitu', 'ndereba', 'murogi', 'kihii',  
          'murutani', 'muiru', 'muici', 'kiratu', 'muruthi', 'thiia',  
          'kimbu', 'nugu', 'kingangi', 'kahiu', 'nyungu', 'gakaraku',  
          'mutune', 'mweru', 'muiru', 'njau', 'mbakuri', 'twara', 'roga',  
          'tura', 'rima', 'enda', 'onja', 'aka', 'endia', 'toga', 'rwara',  
          'ria', 'hokeka', 'uma', 'thoma', 'enyuka', 'ora', 'agana', 'raiha',  
          'kiga', 'kura', 'mwihokeku', 'mwonju', 'muthomu', 'muumu',  
          'nyenyuku', 'njuru', 'njaganu', 'ndaihu', 'ngigu', 'nguru',  
          'inyui', 'ithui', 'nii', 'othee', 'wee', 'atia', 'riria', 'nuu',  
          'ma', 'umuthi', 'niki', 'tene', 'riu', 'hwaiini'], dtype=object)
```

```
[98]: # Viewing target labels  
y = data['Label'].values  
y
```

```
[98]: array([2, 2, 2, 2, 2, 4, 4, 4, 4, 4, 4, 4, 2, 2, 2, 2, 2, 2, 2, 2, 2,  
          2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 4, 4, 4, 4, 4, 4, 4,  
          4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
          0, 3, 3, 3, 3, 3, 3, 3, 3, 1, 1, 1, 1, 1, 1])
```

```
[99]: # Split data into train and test sets  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
          random_state=42)
```

```
[100]: # Tokenize words and pad sequences  
word_to_index = {word: idx + 1 for idx, word in enumerate(set(X))}  
X_train_tokens = np.array([[word_to_index[word] for word in sentence.split()]  
          for sentence in X_train])  
X_test_tokens = np.array([[word_to_index[word] for word in sentence.split()]  
          for sentence in X_test])  
# sequence padding  
max_sequence_length = max(max(len(x) for x in X_train_tokens), max(len(x) for x  
          in X_test_tokens))  
X_train_padded = pad_sequences(X_train_tokens, maxlen=max_sequence_length,  
          padding='post')  
X_test_padded = pad_sequences(X_test_tokens, maxlen=max_sequence_length,  
          padding='post')
```

```
[101]: # Convert labels to one-hot encoding  
num_classes = len(label_encoder.classes_)  
y_train_one_hot = to_categorical(y_train, num_classes=num_classes)  
y_test_one_hot = to_categorical(y_test, num_classes=num_classes)
```

```
[147]: # Defining the BiLSTM model
model = Sequential()
model.add(Embedding(input_dim=len(word_to_index) + 1, output_dim=100,
    ↪input_length=max_sequence_length))
model.add(Bidirectional(LSTM(32, return_sequences=True)))
#Flatten layer to match the output shape
model.add(Flatten())
model.add(Dense(num_classes, activation='softmax'))
```

```
[148]: # model summary
model.summary()
```

Model: "sequential_16"

Layer (type)	Output Shape	Param #
embedding_16 (Embedding)	(None, 1, 100)	8100
bidirectional_16 (Bidirectional)	(None, 1, 64)	34048
flatten_13 (Flatten)	(None, 64)	0
dense_16 (Dense)	(None, 5)	325

=====
 Total params: 42473 (165.91 KB)
 Trainable params: 42473 (165.91 KB)
 Non-trainable params: 0 (0.00 Byte)
 =====

```
[149]: # Defining loss function and the appropriate optimizer
model.compile(optimizer='adam', loss='categorical_crossentropy',
    ↪metrics=['accuracy'])
```

```
[150]: # Training the model
model.fit(X_train_padded, y_train_one_hot,
    epochs=10, batch_size=16,
    validation_data=(X_test_padded,
        y_test_one_hot))
```

Epoch 1/10
 4/4 [=====] - 5s 313ms/step - loss: 1.6078 - accuracy: 0.3125 - val_loss: 1.6063 - val_accuracy: 0.2941
 Epoch 2/10
 4/4 [=====] - 0s 19ms/step - loss: 1.5965 - accuracy: 0.5469 - val_loss: 1.6024 - val_accuracy: 0.4118

```

Epoch 3/10
4/4 [=====] - 0s 22ms/step - loss: 1.5857 - accuracy:
0.6875 - val_loss: 1.5983 - val_accuracy: 0.4706
Epoch 4/10
4/4 [=====] - 0s 16ms/step - loss: 1.5737 - accuracy:
0.6875 - val_loss: 1.5951 - val_accuracy: 0.4118
Epoch 5/10
4/4 [=====] - 0s 19ms/step - loss: 1.5614 - accuracy:
0.7188 - val_loss: 1.5913 - val_accuracy: 0.3529
Epoch 6/10
4/4 [=====] - 0s 16ms/step - loss: 1.5468 - accuracy:
0.7188 - val_loss: 1.5879 - val_accuracy: 0.2941
Epoch 7/10
4/4 [=====] - 0s 20ms/step - loss: 1.5299 - accuracy:
0.7344 - val_loss: 1.5844 - val_accuracy: 0.2941
Epoch 8/10
4/4 [=====] - 0s 20ms/step - loss: 1.5109 - accuracy:
0.7344 - val_loss: 1.5805 - val_accuracy: 0.2941
Epoch 9/10
4/4 [=====] - 0s 15ms/step - loss: 1.4864 - accuracy:
0.7500 - val_loss: 1.5764 - val_accuracy: 0.2941
Epoch 10/10
4/4 [=====] - 0s 15ms/step - loss: 1.4591 - accuracy:
0.8125 - val_loss: 1.5721 - val_accuracy: 0.2941

```

[150]: <keras.src.callbacks.History at 0x78f652b2faf0>

```

[151]: # Evaluating the model
loss, accuracy = model.evaluate(X_test_padded, y_test_one_hot)
print(f'Test Accuracy: {accuracy * 100:.2f}%')

```

```

1/1 [=====] - 0s 33ms/step - loss: 1.5721 - accuracy:
0.2941
Test Accuracy: 29.41%

```

```

[143]: # Function to Test of the model can tag a word to its respective POS
def predict_pos_tags(model, word_to_index, label_encoder):
    while True:
        # User to enter words
        user_input = input("Enter a sentence or a list of words (type 'exit' to
quit): ")
        if user_input.lower() == 'exit':
            break

        # Tokenize and pad the input
        tokens = [word_to_index[word] for word in user_input.split()]

```

```

        padded_tokens = pad_sequences([tokens], maxlen=max_sequence_length,
↪padding='post')

        # Predict POS tags
        predictions = model.predict(padded_tokens)

        # Decode predicted labels
        predicted_labels = label_encoder.inverse_transform(np.
↪argmax(predictions, axis=1))

        # Print the words along with their predicted POS tags
        for word, pos_tag in zip(user_input.split(), predicted_labels):
            print(f"{word}: {pos_tag}")

# Call the predict_pos_tags function
predict_pos_tags(model, word_to_index, label_encoder)

```

```

Enter a sentence or a list of words (type 'exit' to quit): ihua
1/1 [=====] - 1s 757ms/step
ihua: Noun
Enter a sentence or a list of words (type 'exit' to quit): muaki
1/1 [=====] - 0s 30ms/step
muaki: Noun
Enter a sentence or a list of words (type 'exit' to quit): rehe
1/1 [=====] - 0s 32ms/step
rehe: Verb
Enter a sentence or a list of words (type 'exit' to quit): njaganu
1/1 [=====] - 0s 21ms/step
njaganu: Verb
Enter a sentence or a list of words (type 'exit' to quit): nii
1/1 [=====] - 0s 21ms/step
nii: Verb
Enter a sentence or a list of words (type 'exit' to quit): mwihokeku
1/1 [=====] - 0s 23ms/step
mwihokeku: Verb
Enter a sentence or a list of words (type 'exit' to quit): exit

```

The Poor Performance of the model is caused by having few instances in training set ,thus the model is not able to learn alot of context on these words,therefore it will perform very dismally on new data